

# Review on Lec 1 & Lec 2

## Hello CUDA

```

__global__ void addKernel(int * const a, const int * const b, const int *
const c)
{
    const unsigned int i = threadIdx.x;
    c[i] = a[i] + b[i];
}

void main(){
    .....
    int *dev_a,*dev_b,*dev_c;
    cudaMalloc((void**)&dev_c, 128* sizeof(int));
    .....
    cudaMemcpy(dev_a, a, 128* sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, 128* sizeof(int), cudaMemcpyHostToDevice);

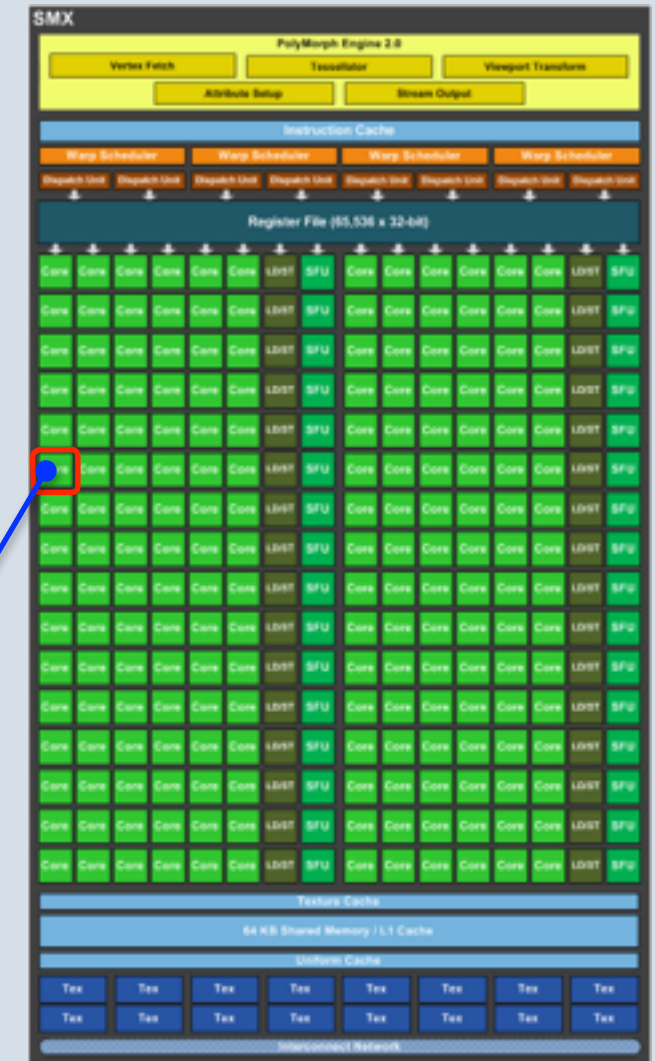

    // Launch a kernel on the GPU with one thread for each element.
    addKernel<<<1, 128>>>>(dev_c, dev_a, dev_b);

    cudaMemcpy(c, dev_c, 128* sizeof(int), cudaMemcpyDeviceToHost);

    cudaFree(dev_c);
    .....
}

```

## Kepler SMX



# **Lec 3 CUDA Software Abstraction**

**Dong Li, Tonghua Su**  
School of Software  
Harbin Institute of Technology

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# Outline


- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# Design Philosophy

**CPU: Latency Oriented Cores**

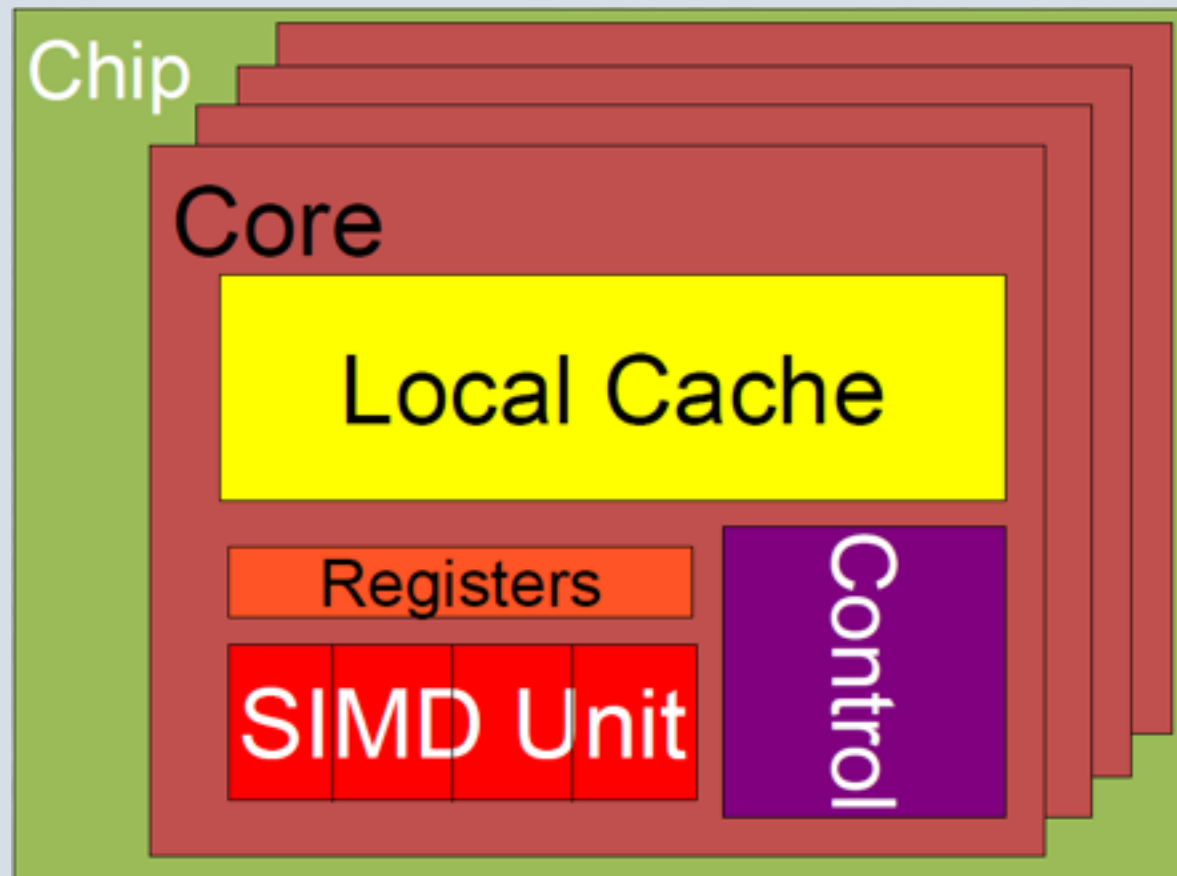


**GPU: Throughput Oriented Cores**

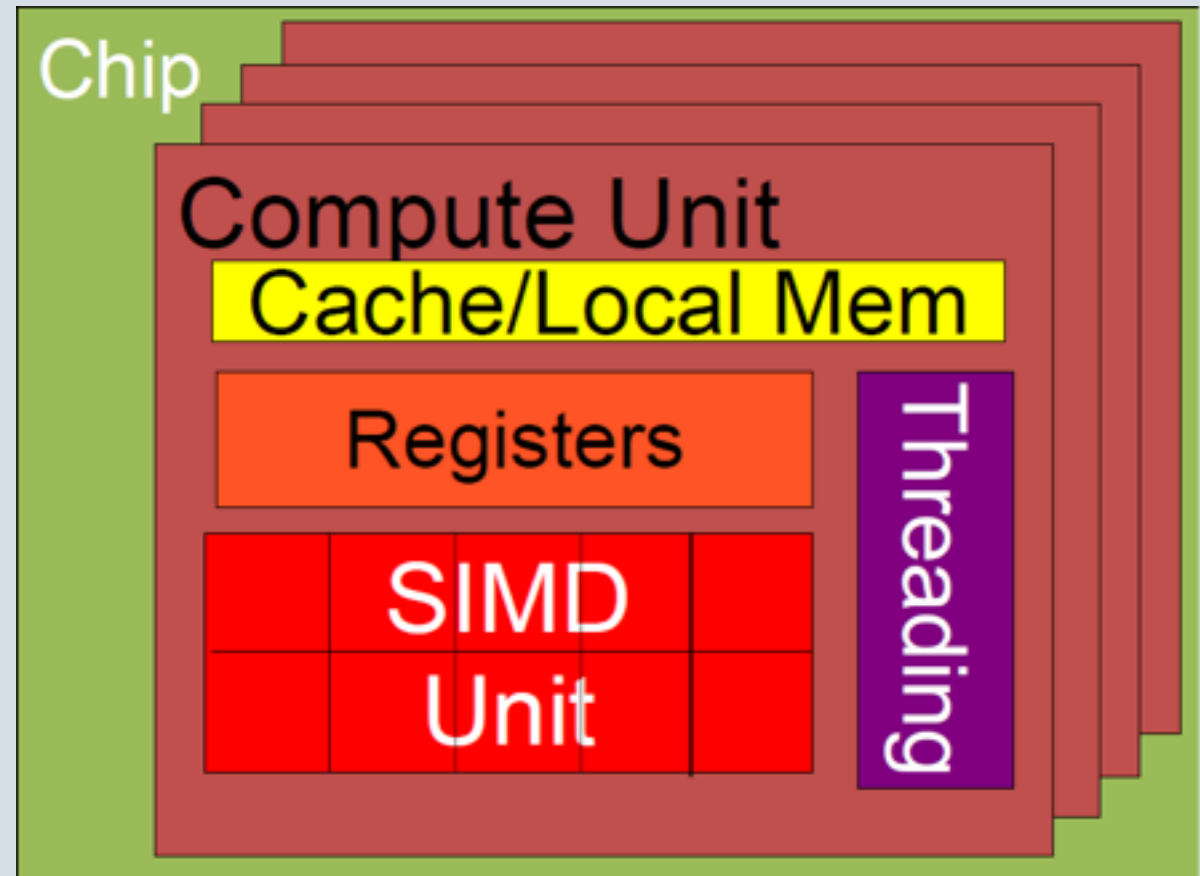


# Design Philosophy

## CPU: Latency Oriented Cores



## GPU: Throughput Oriented Cores



# Multithreading

## ●SIMD in CPU

- ✓ All cores execute the same instructions simultaneously, but with different data
- ✓ Similar to vector computing on CRAY supercomputers
- ✓ e. g. SSE4, AVX instruction set

## ●SIMT in SMX

- ✓ Multithreaded CUDA core
- ✓ Threads on each SMX execute in group sharing same instruction
- ✓ Fine-grained parallelism
- ✓ Natural for graphics processing and much scientific computing
- ✓ SIMT is also a natural choice for many-core chips to simplify each core

# Multithreading

- **Thread: instruction stream with own PC and data**
  - ✓ **Owning private register, private memory, program counter and thread execution state**
  - ✓ **Thread Level Parallelism(TLP): Exploit the parallelism inherent between threads**
- **Multithreading**
  - ✓ **Multiple threads to share the functional units of 1 processor via overlapping**
  - ✓ **Processor must duplicate independent state of each thread**
    - e.g., a separate copy of register file, a separate PC
  - ✓ **Often, hardware for fast thread switch**
  - ✓ **Memory to be shared**
  - ✓ **Solving the memory access stall**



# GPU Multithreading

## ●Multithreaded Hardware

- ✓ **CUDA core is multithreaded processor**

  - Supporting 96 threads in GTX 8800

- ✓ **Run in group of 32 threads (called a warp)**

- ✓ **Zero-cost “context switching”**

  - Each thread has its own registers ( which limits the number of active threads

- ✓ **Shared memory/L1 cache**

## ●Fine-grained multithreading

- ✓ **Able to switch between warps on each instruction**

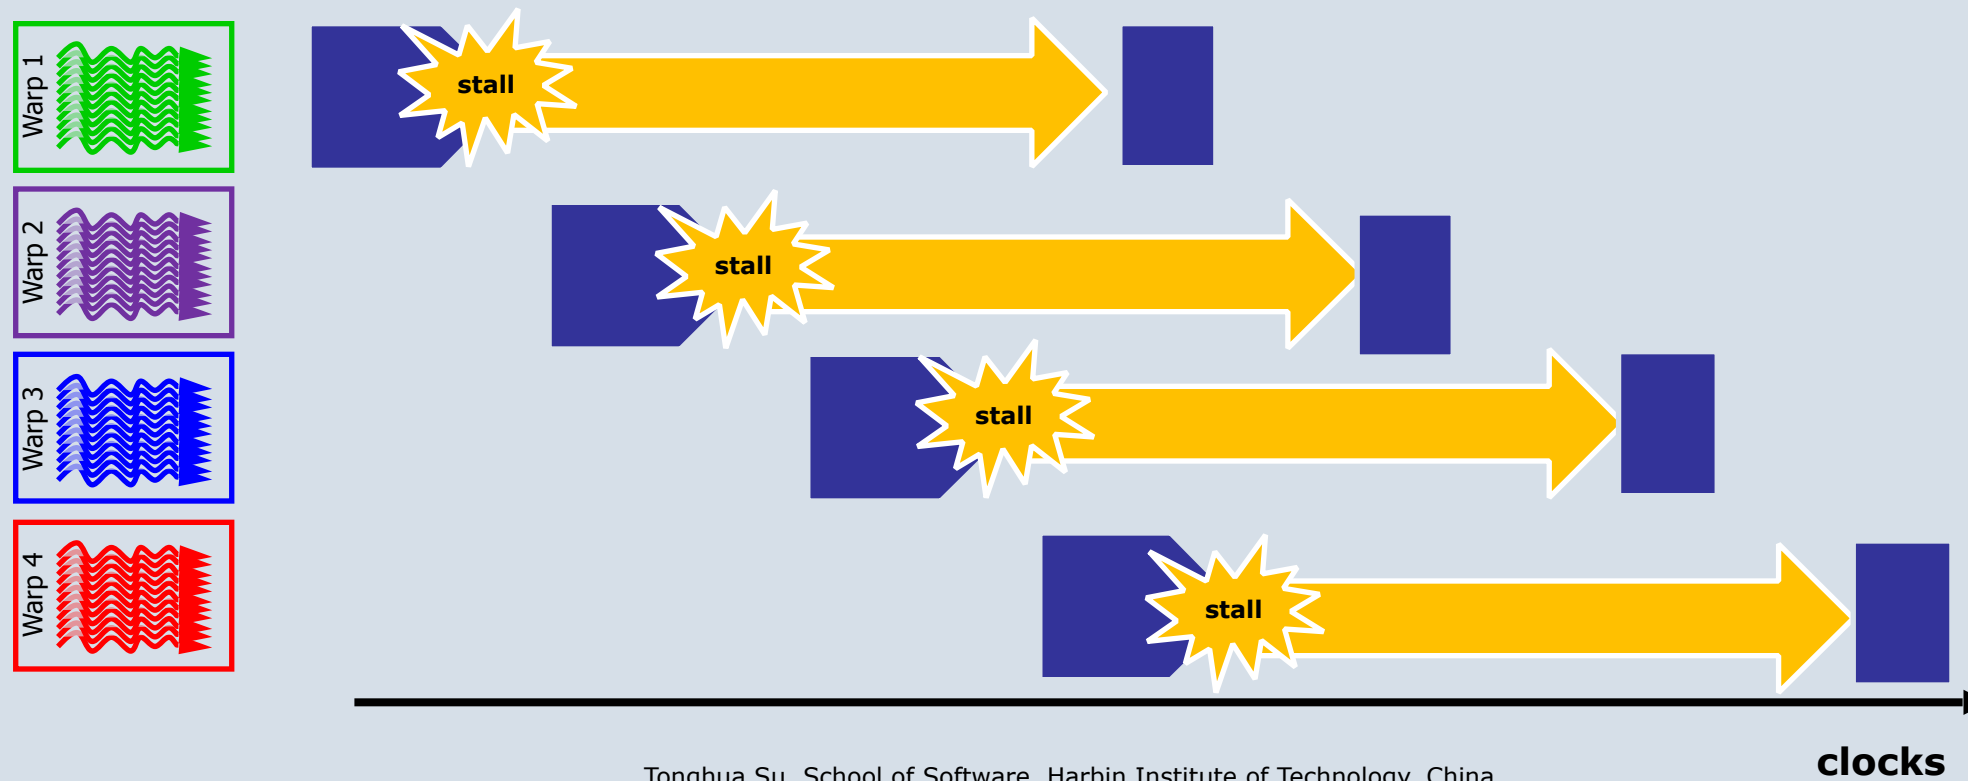
- ✓ **Schedule without pre-emption**

- ✓ **Scheduling eligible warps in a round-robin fashion**

# GPU Multithreading

## ●Hiding Latency Stalls

- ✓ Fetch → Decode → Execute → **Memory** → Writeback
- ✓ Execution alternates between “active” warps, with warps becoming temporarily “inactive” when waiting for data
- ✓ Lots of active warps is the key to high performance



# Quiz

- 定量计算内存访问带来的停滞
  - ✓ 假设从显存读取数据的延时是 **400**时钟周期
  - ✓ 如果每个线程束 (**warp**) 可以运行**10**个周期，那么需要多少个就绪的线程束才能掩盖停滞带来的时间缝隙?

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# CUDA

- **CUDA(Compute Unified Device Architecture) is developed by Nvidia around 2007**
  - ✓ 2-4 week learning curve for those with experience of OpenMP and MPI programming
  - ✓ large user community on NVIDIA forums
- **CUDA is a parallel computing platform and programming model that enables dramatic increases in computing performance by harnessing the power of the GPU**

# CUDA

## ●CUDA as Parallel Computing Platform

- ✓ **Language:** CUDA C which based on C with some extensions(extensive C++ support)
- ✓ **Editor:** Eclipse/Visual Studio
- ✓ **Compiler:** nvcc
- ✓ **SDK:** CUDA toolkit, Libraries, Samples
- ✓ **Profiler & Debugger:** Nsight

## ●CUDA as Programming Model

- ✓ **Software Abstraction of GPU hardwares**
- ✓ **Independent to OSs, CPUs, Nvidia GPUs**


# CUDA Abstraction

## ●CUDA Virtualizes the Physical Hardware

- ✓ thread is a virtualized CUDA cores (registers, PC, state)
- ✓ block is a virtualized streaming multiprocessor (threads, shared mem.)

## ●Scheduled onto Physical Hardware without Pre-emption

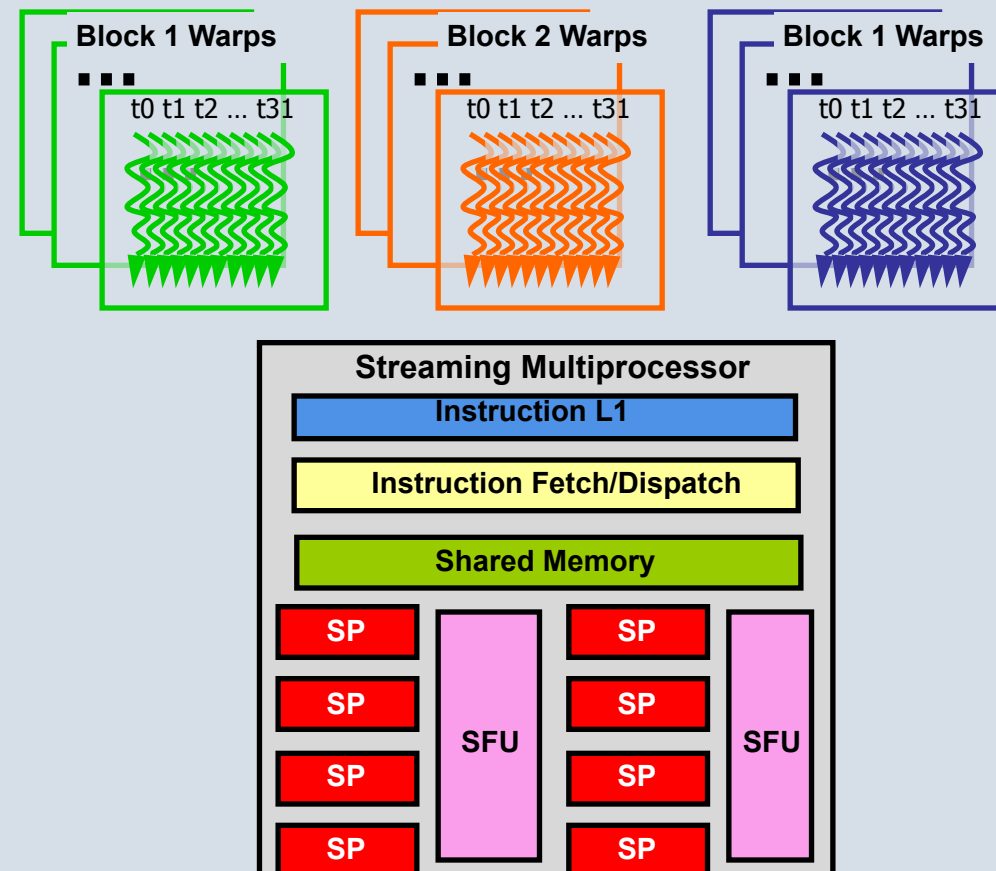
- ✓ threads/blocks launch & run to completion/suspension
- ✓ blocks should be independent



# CUDA Abstraction

## ●Key Parallel Abstractions in CUDA

- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ Lightweight synchronization primitives

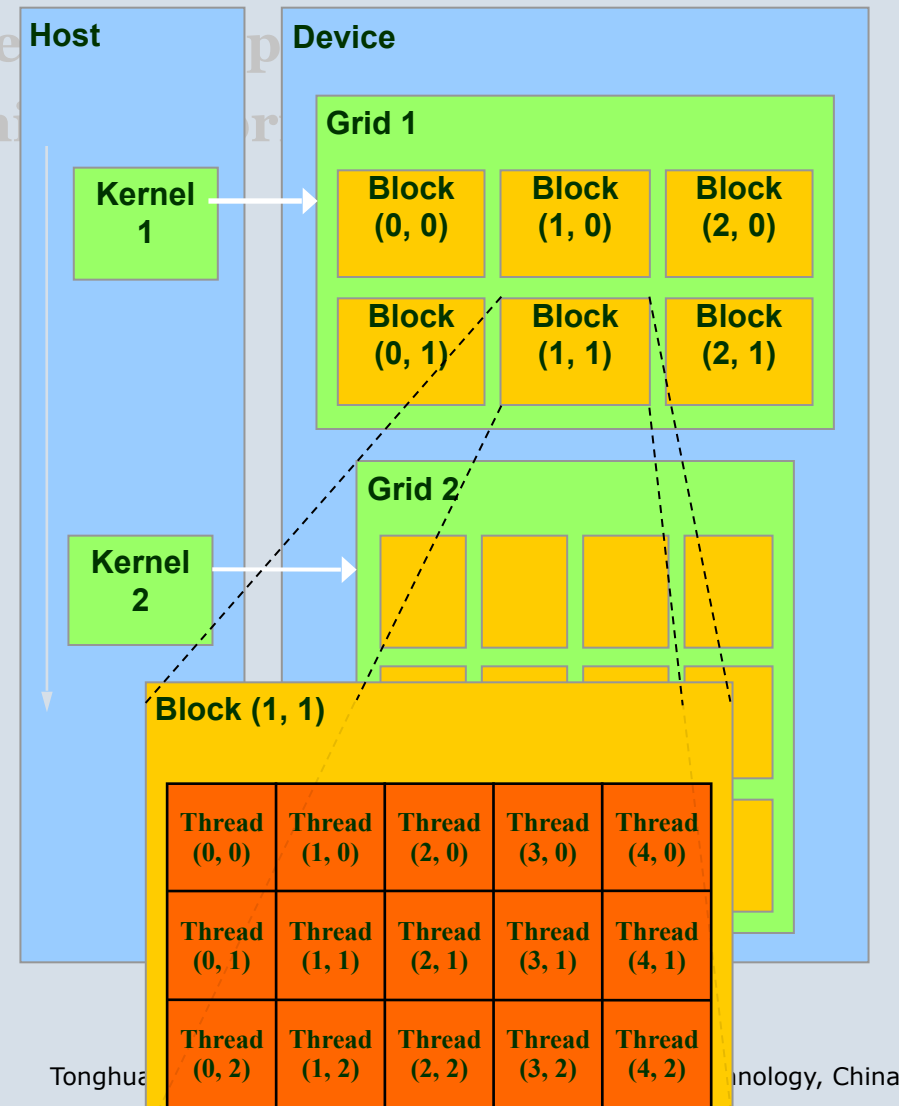




# CUDA Abstraction

## ●Key Parallel Abstractions in CUDA


- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model
- ✓ Lightweight synchronization




# Thread Hierarchy

## ● Thread → Block → Grid

```
void main(){  
    .....  
    int *dev_a,*dev_b,*dev_c;  
    .....  
    addKernel<<<1, 128>>>(dev_c, dev_a, dev_b);  
    .....  
}
```




```
void main(){  
    .....  
    int *dev_a,*dev_b,*dev_c;  
    .....  
    addKernel<<<100, 128>>>(dev_c, dev_a, dev_b);  
    .....  
}
```

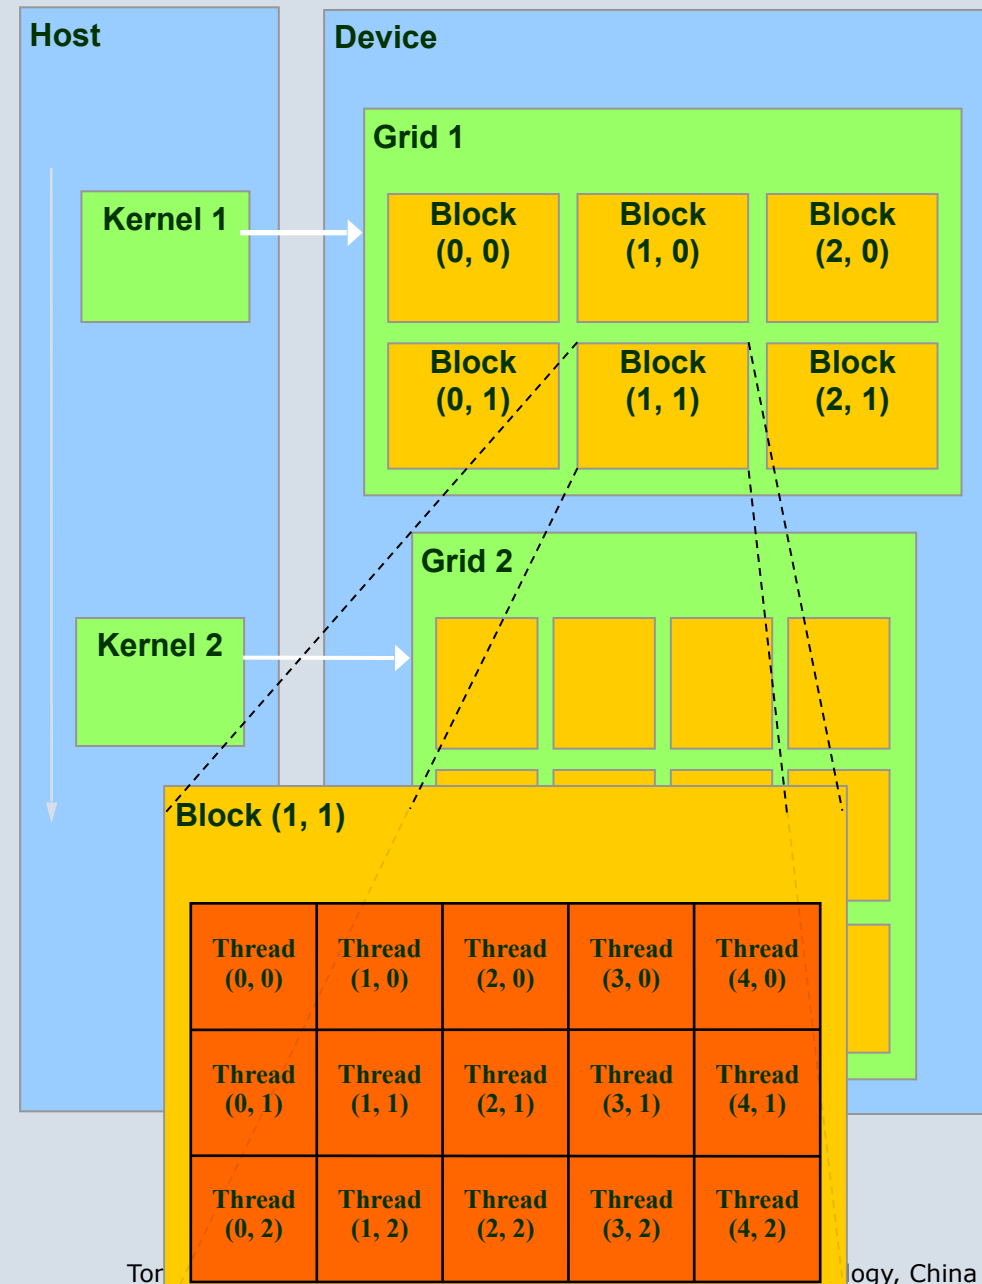


# Thread Hierarchy

## ● Thread Mapping



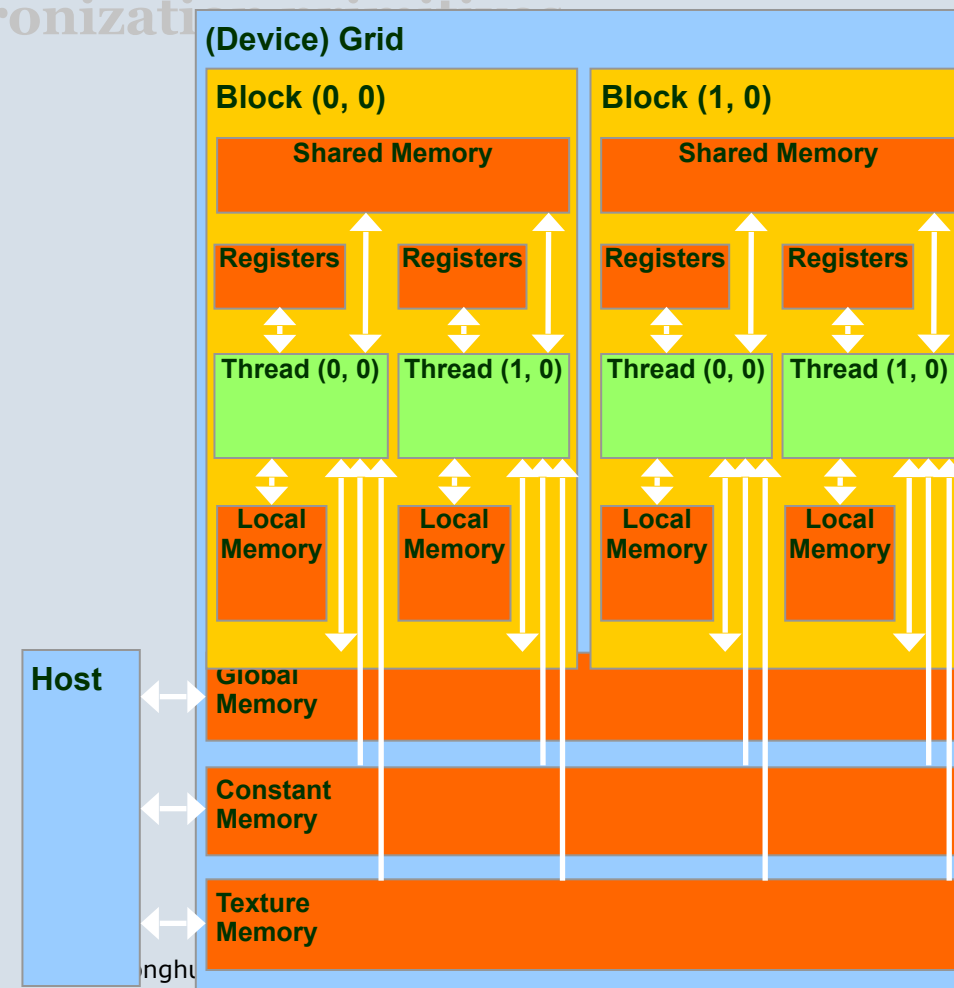
# Thread Hierarchy



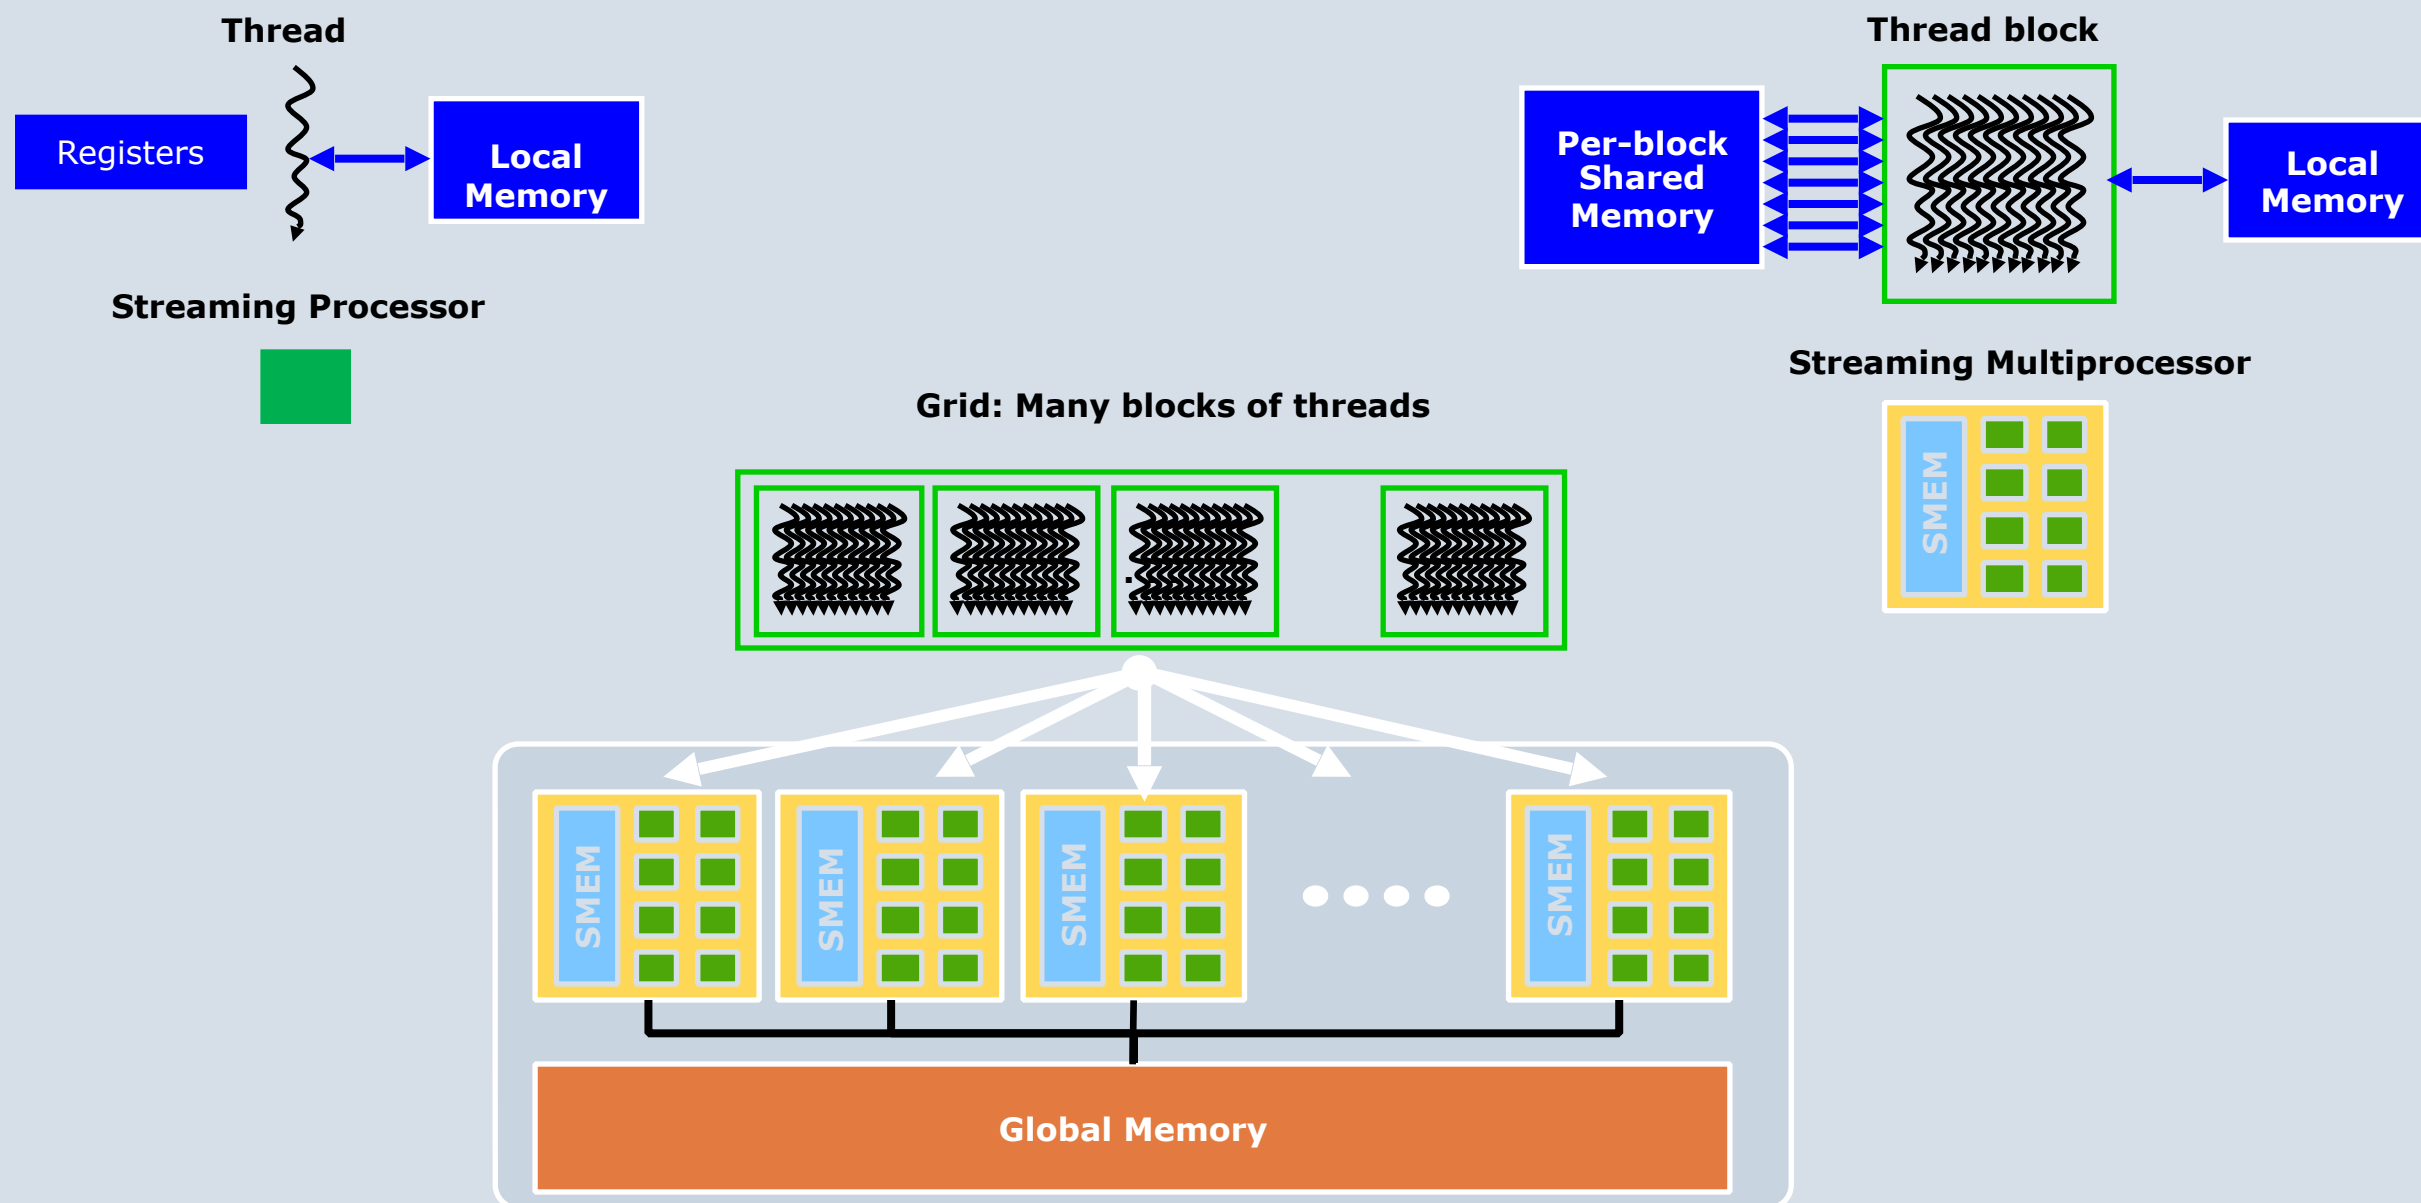
# CUDA Abstraction

## ●Key Parallel Abstractions in CUDA

- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ Lightweight synchronization primitives

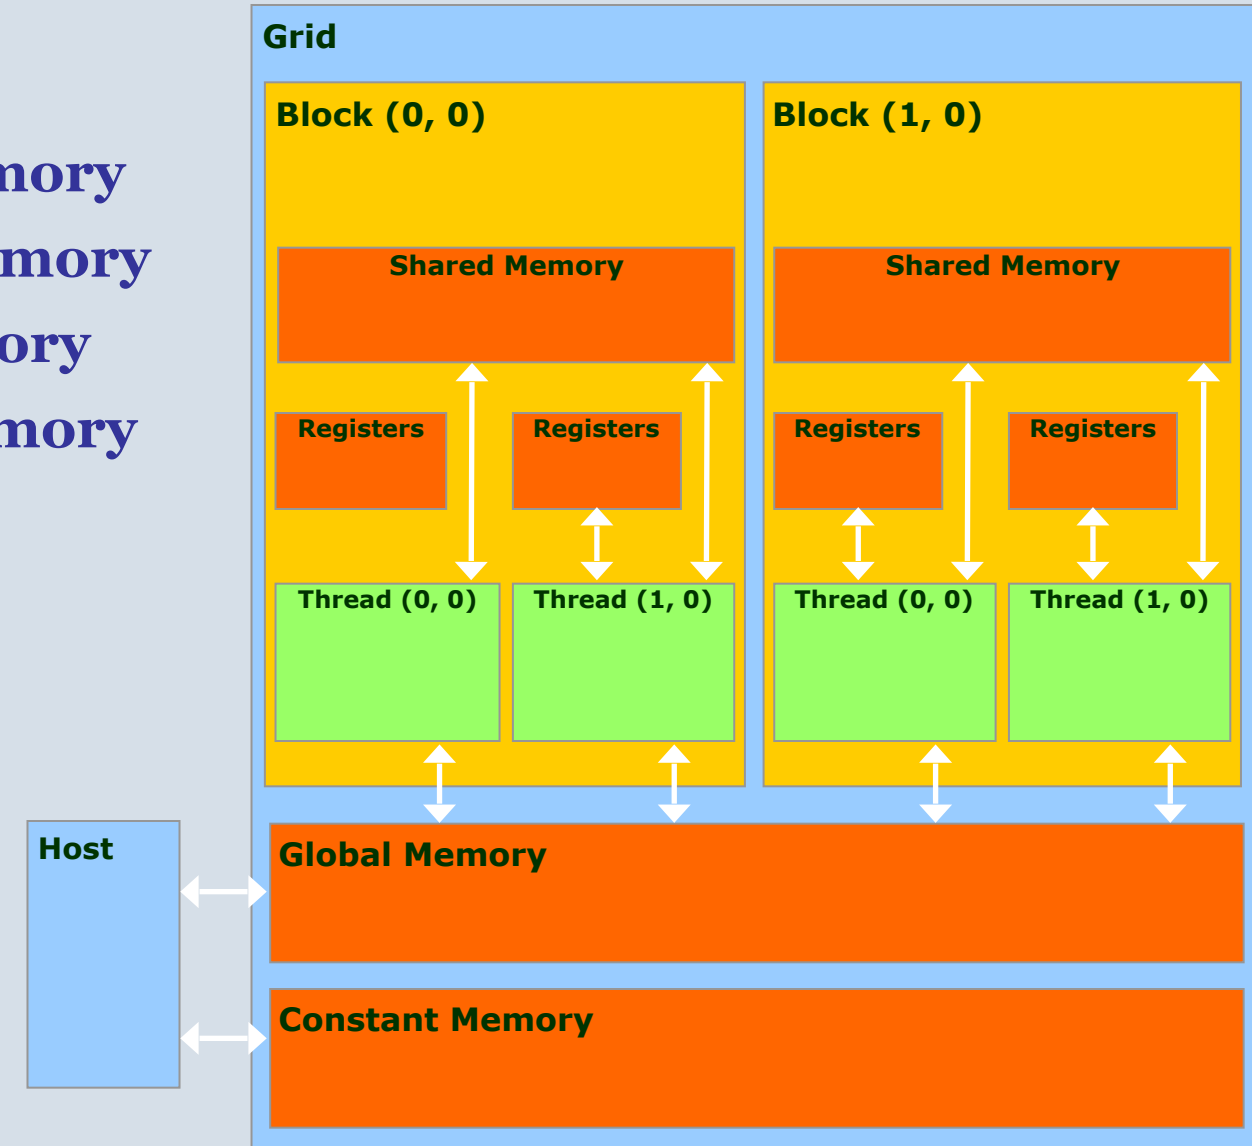


# Memory Model



# Memory Model

- Each thread can:
  - ✓ Read/write per-thread **registers**
  - ✓ Read/write per-thread **local memory**
  - ✓ Read/write per-block **shared memory**
  - ✓ Read/write per-grid **global memory**
  - ✓ Read/only per-grid **constant memory**



# CUDA Abstraction

## ●Key Parallel Abstractions in CUDA

- ✓ Hierarchy of concurrent threads
- ✓ Shared memory model for cooperating threads
- ✓ **Lightweight synchronization primitives**



# Synchronization

## ●Global Synchronization

- ✓ Finish a kernel and start a new one
- ✓ All writes from all threads complete before a kernel finishes

```
step1<<<grid1,blk1>>>(...);  
// The system ensures that all writes from step1  
complete.  
step2<<<grid2,blk2>>>(...);
```

- ✓ Would need to decompose kernels into before and after parts

# Synchronization

## ●Threads Synchronization

- ✓ To ensure the threads visit the shared memory in order
- ✓ `__syncthreads()`

```
__global__ void adj_diff(int *result, int *input)
{
    int tx = threadIdx.x;
    // allocate a __shared__ array, one element per thread
    __shared__ int s_data[BLOCK_SIZE];
    // each thread reads one element to s_data
    unsigned int i = blockDim.x * blockIdx.x + tx;
    s_data[tx] = input[i];
    // avoid race condition: ensure all loads complete before continuing
    __syncthreads();
    if(tx > 0) result[i] = s_data[tx] - s_data[tx-1];
    else if(i > 0)
    { // handle thread block boundary
        result[i] = s_data[tx] - input[i-1];
    }
}
```

# Synchronization

## ●Race Conditions

- ✓ What is the value of `a` in thread 0?
- ✓ What is the value of `a` in thread 127?

threadId:0	threadId:127
<code>// vector[0] was equal to 0</code>	
<code>vector[0] += 5;</code>	<code>vector[0] += 1;</code>
<code>...</code>	<code>...</code>
<code>a = vector[0];</code>	<code>a = vector[0];</code>

- ✓ CUDA provides **atomic** operations to deal with this problem

# Synchronization

## ●Atomics

- ✓ An atomic operation guarantees that only a single thread has access to a piece of memory while an operation completes
- ✓ Different types of atomic instructions:
  - `atomic{Add, Sub, Exch, Min, Max, Inc, Dec, CAS, And, Or, Xor}`
- ✓ Atomics are slower than normal load/store
- ✓ You can have the whole machine queuing on a single location in memory
- ✓ More types in Fermi
- ✓ Atomics unavailable on G80!

# Synchronization

## ●Atomics

```
// Determine frequency of colors in a picture
// colors have already been converted into ints
// Each thread looks at one pixel and increments
// a counter atomically
__global__ void histogram(int* color,
                          int* buckets)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    int c = colors[i];
    atomicAdd(&buckets[c], 1);
}
```

# Quiz

## ● 填空

- ✓ CUDA中，( )是对**CUDA core**的抽象，( )是对**SMX**的抽象
- ✓ CUDA的三大抽象技术分别是( )、( )和( )

## ● 简答

- ✓ 请自查资料，说明**atomicCAS**的作用，如可能给出实例

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# CUDA Function

## ●Kernel Function

- ✓ C-style function that is executed in parallel by more than one thread
- ✓ defined using **\_\_global\_\_** declaration specifier and returns **void**

```
__global__ void addKernel(int * const a, const int * const b, const int * const c)
{
    const unsigned int i = threadIdx.x;
    c[i] = a[i] + b[i];
}
```

## ●CUDA Function

- ✓ **\_\_device\_\_** and **\_\_host\_\_** can be used together

<i>Syntax</i>	<i>Executed on</i>	<i>Only callable from</i>
<b>__device__</b> float DeviceFunc()	<b>device</b>	<b>device</b>
<b>__global__</b> void KernelFunc()	<b>device</b>	<b>host, device(since sm3.x)</b>
<b>__host__</b> float HostFunc()	<b>host</b>	<b>host</b>



# Kernel Function

## ●Kernel Invocation Looks Like:

**KernelFunc<<<gridDim, blockDim, nSMem,iStream>>>(args);**

- ✓ **<<<>>>** identify the execution configuration of the kernel
- ✓ **gridDim** is the dimension of the grid
- ✓ **blockDim** is the dimension of each block
- ✓ **args** is a limited number of arguments, usually mainly pointers to arrays in graphics memory, and some constants which get copied by value
- ✓ The more general form allows gridDim and blockDim to be 2D or 3D (**dim3** type) to simplify application programs

# Kernel Function

## ●Kernel Dimensional Limitation:

```
C:\ProgramData\NVIDIA Corporation\CUDA Samples\v6.5\bin\win64\Release>deviceQuery.exe
deviceQuery.exe Starting...


  CUDA Device Query (Runtime API) version (CUDA RT static linking)

Detected 1 CUDA Capable device(s)

Device 0: "GeForce GTX 680"
  CUDA Driver Version / Runtime Version      6.5 / 6.5
  CUDA Capability Major/Minor version number: 3.0
  Total amount of global memory:              2048 MBytes (2147483648 bytes)
  ( 8) Multiprocessors, (192) CUDA Cores/MP:  1536 CUDA Cores
  GPU Clock rate:                            1059 MHz (1.06 GHz)
  Memory Clock rate:                          3004 Mhz
  Memory Bus Width:                           256-bit
  L2 Cache Size:                              524288 bytes
  Maximum Texture Dimension Size (x,y,z)      1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
  Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
  Total amount of constant memory:             65536 bytes
  Total amount of shared memory per block:     49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                  32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:         1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                           512 bytes
  Concurrent copy and kernel execution:        Yes with 1 copy engine(s)
  Run time limit on kernels:                    No
  Integrated GPU sharing Host Memory:           No
  Support host page-locked memory mapping:      Yes
  Alignment requirement for Surfaces:           Yes
  Device has ECC support:                       Disabled
```

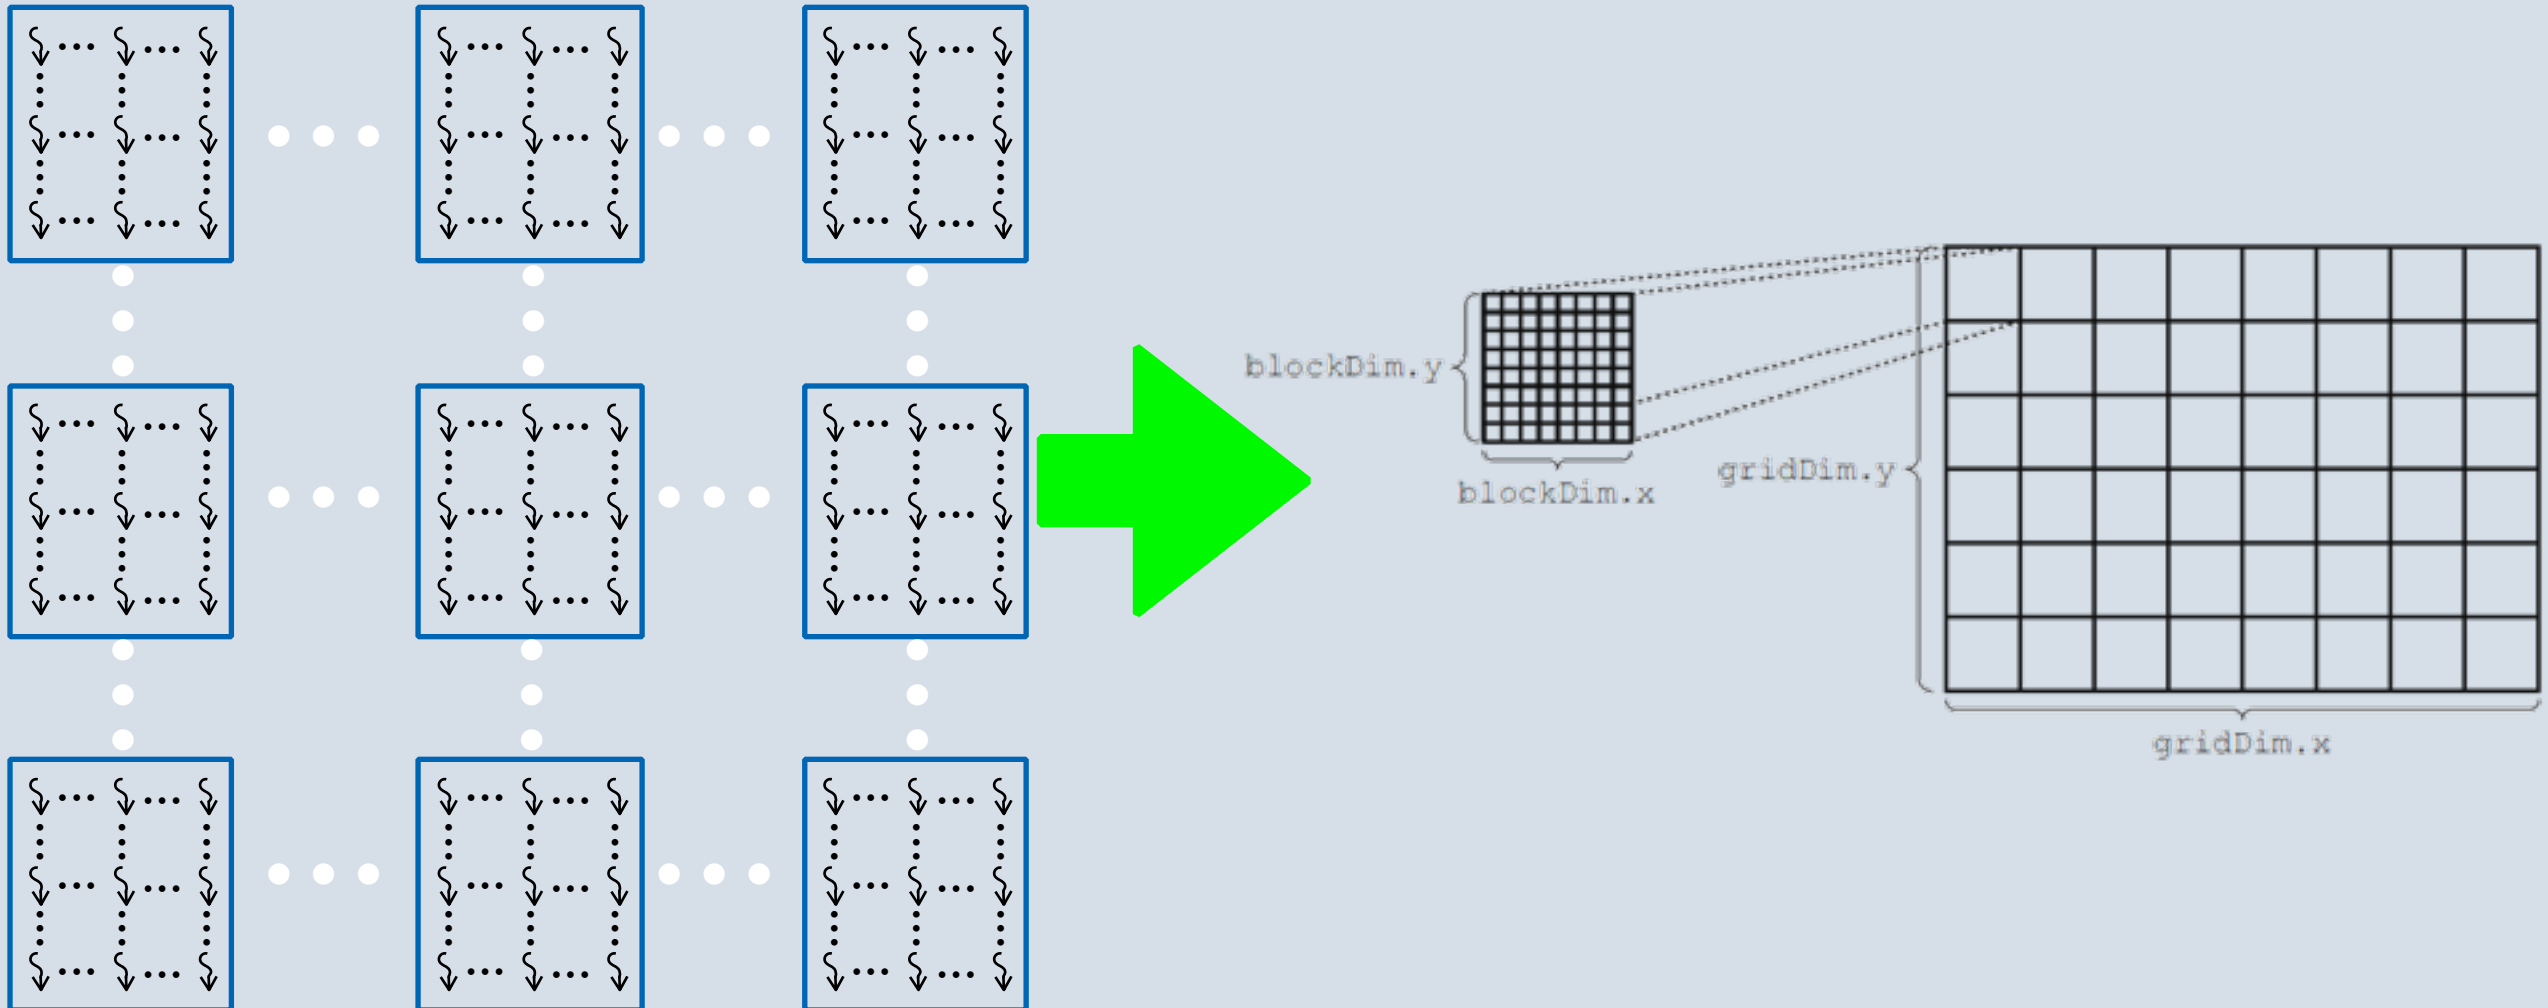
# Thread Indexing

## ● 1D block and 1D grid




# Thread Indexing

## ● 2D block and 2D grid



# Thread Indexing


## ● 2D block and 2D grid




# Data Indexing

- **2D block and 2D grid**

- ✓ Using the blocks to tile the data




**16\*16 block**



# Thread Indexing

## ● 3D block and 3D grid



# Occupancy

- **Occupancy is a useful metric**

- ✓ **the ratio of the number of active warps per multiprocessor to the maximum number of warps that can be active on the multiprocessor at once**
- ✓ **the denominator is a constant that depends only on compute capability of the device**
- ✓ **the numerator of this expression is a function of the following:**
  - compute capability
  - threads per block
  - registers per thread
  - shared memory per block and shared memory configuration



# Quiz 1

- 计算下面内核配置在**Kepler**架构GPU上的设备占用率

```
dim3 blockDim, gridDim;  
blockDim.x = 16;  
blockDim.y = 16;  
gridDim.x = (1000-blockDim.x+1)/blockDim.x;  
gridDim.y = (3755-blockDim.y+1)/blockDim.y;  
distKernel<<<gridDim,blockDim>>>(...)
```

## Quiz 2

- 计算线程对应的数组下标

- ✓ 灰度图像大小为100\*100大小，每个元素的取值在0~255之间，
- ✓ 现在按行展开的方式存储在一维数组**a**里
- ✓ 现在启动如下配置的内核，每个线程操作对应于**a**的一个元素，请问第**(p,q)**个**block**的第**(m,n)**个线程要操作数组**a**的第几个元素？

```
blockDim.x = 16;  
blockDim.y = 16;  
gridDim.x = (100-blockDim.x+1)/blockDim.x;  
gridDim.y = (100-blockDim.y+1)/blockDim.y;  
calKernel<<<gridDim,blockDim>>>(...)
```

# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# CUDA Scheduling

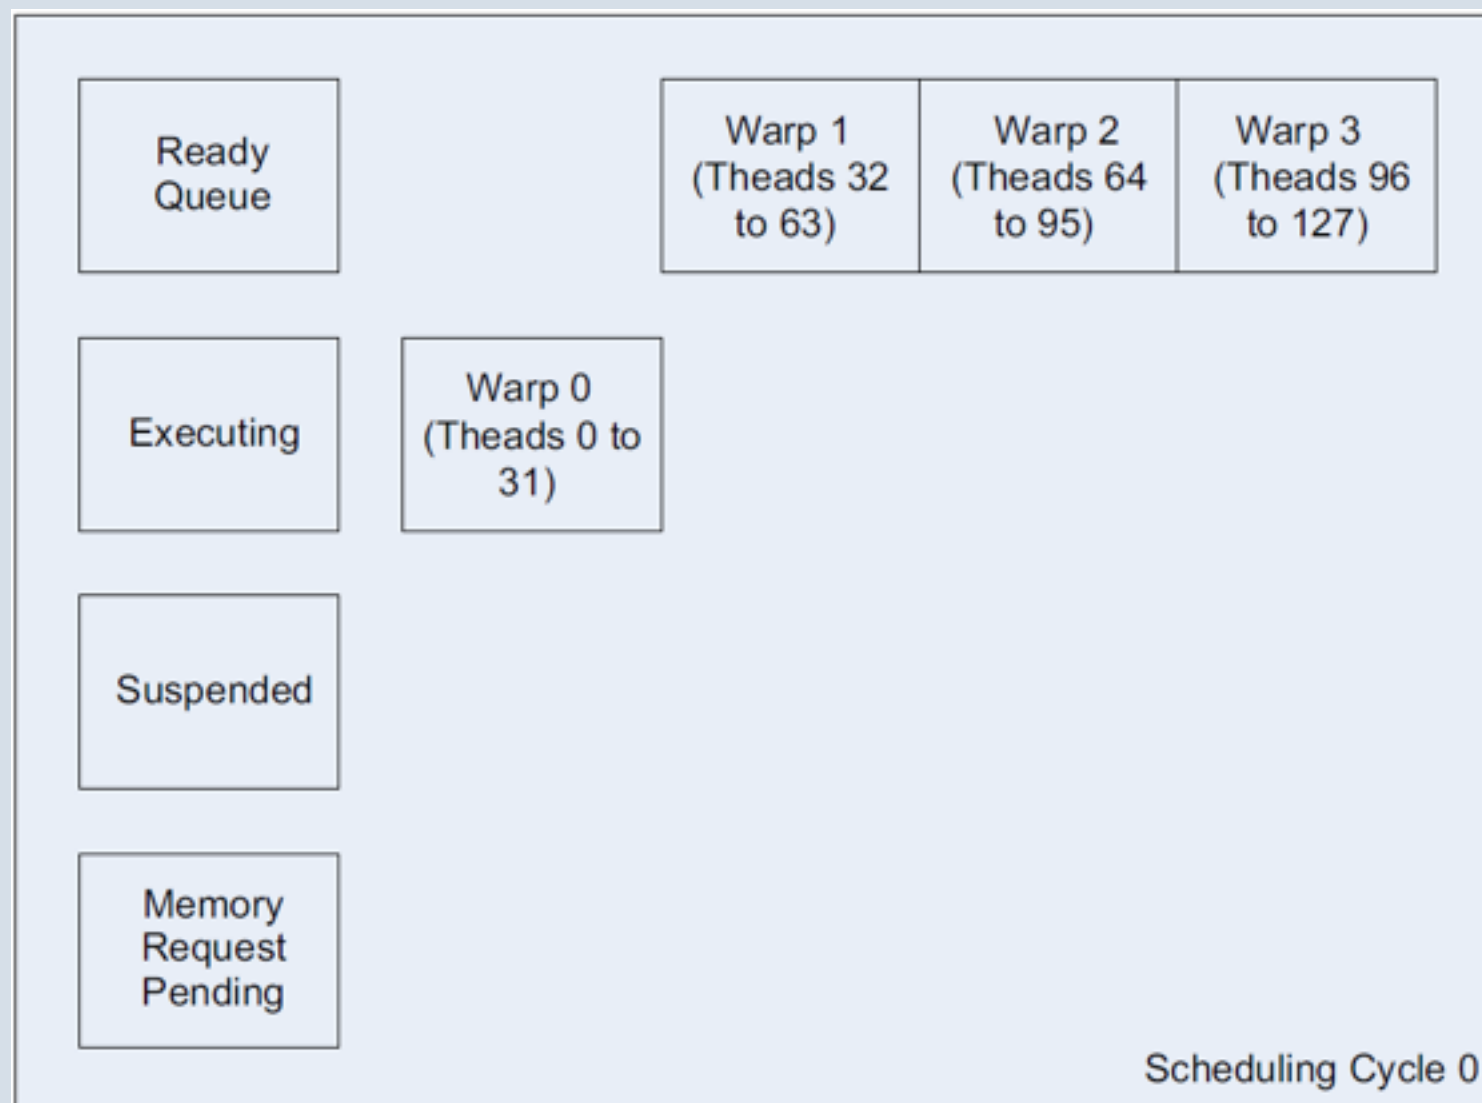
- **At a lower level, within the GPU:**

- ✓ each block of the execution kernel executes on an SMX
- ✓ if the number of blocks exceeds the number of SMXs, then more than one will run at a time on each SMX if there are enough registers and shared memory, and the others will wait in a queue and execute later
- ✓ all threads within one block can access local shared memory but can't see what the other block are doing (even if they are on the same SMX)
- ✓ there are no guarantees on the order in which the blocks execute

# CUDA Scheduling

## ●Block Scheduling

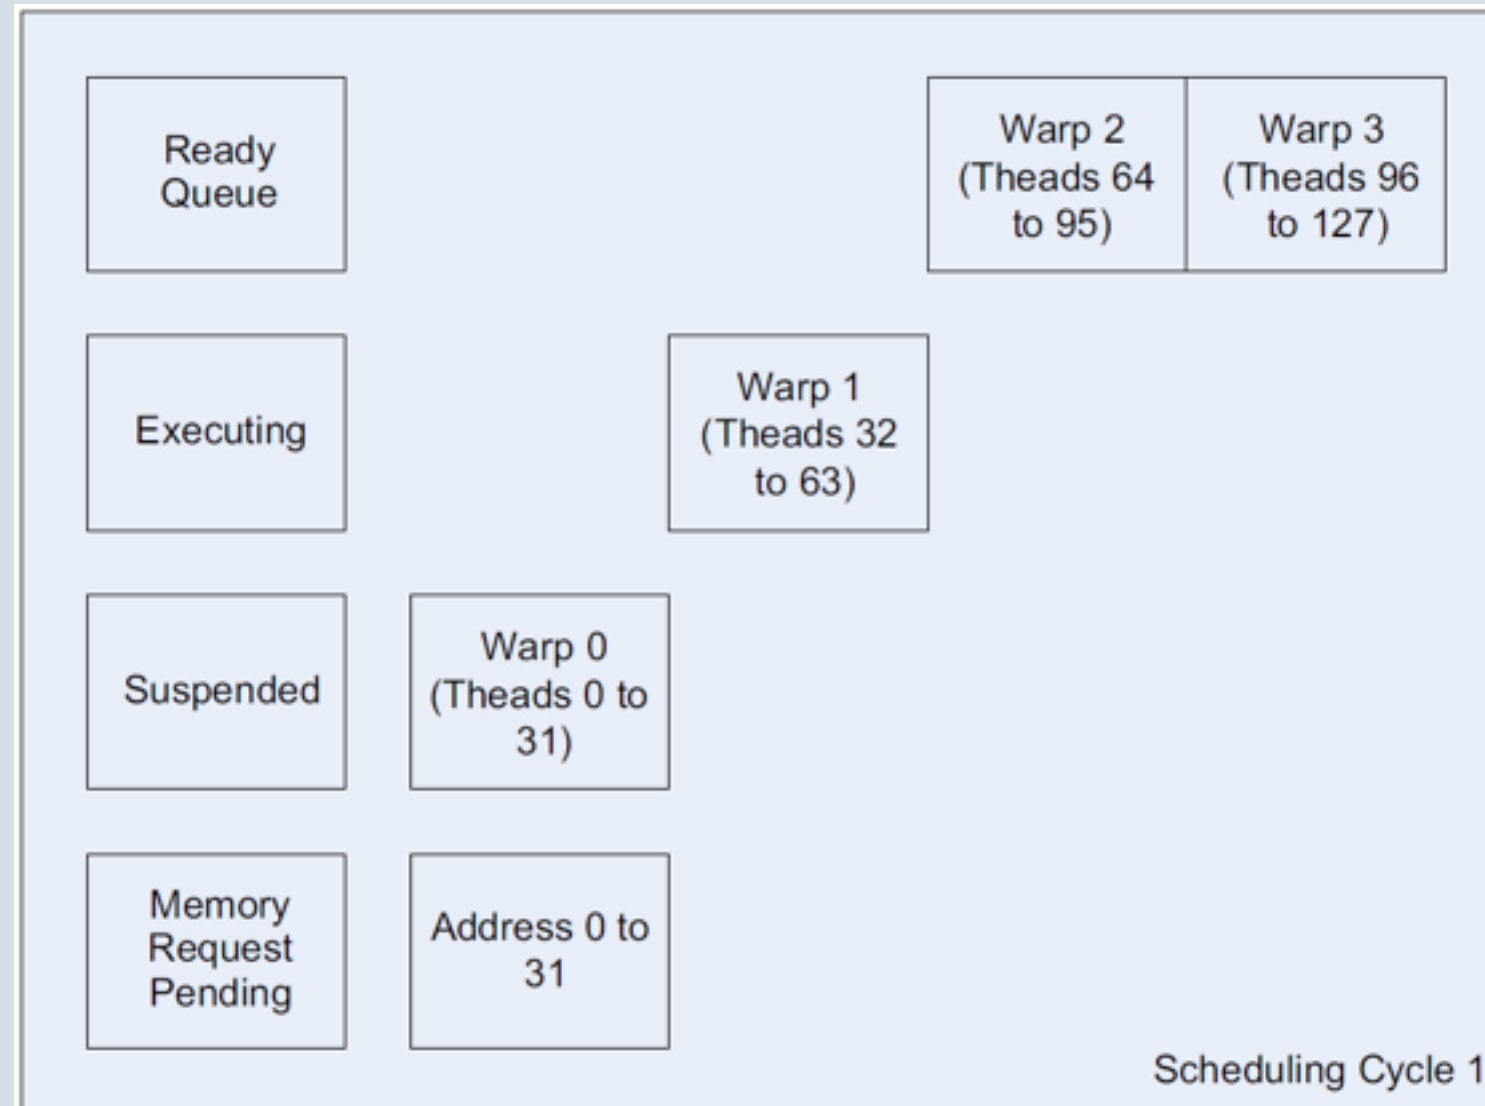
- ✓ Execute in warps of 32 threads



# CUDA Scheduling

## ●Block Scheduling

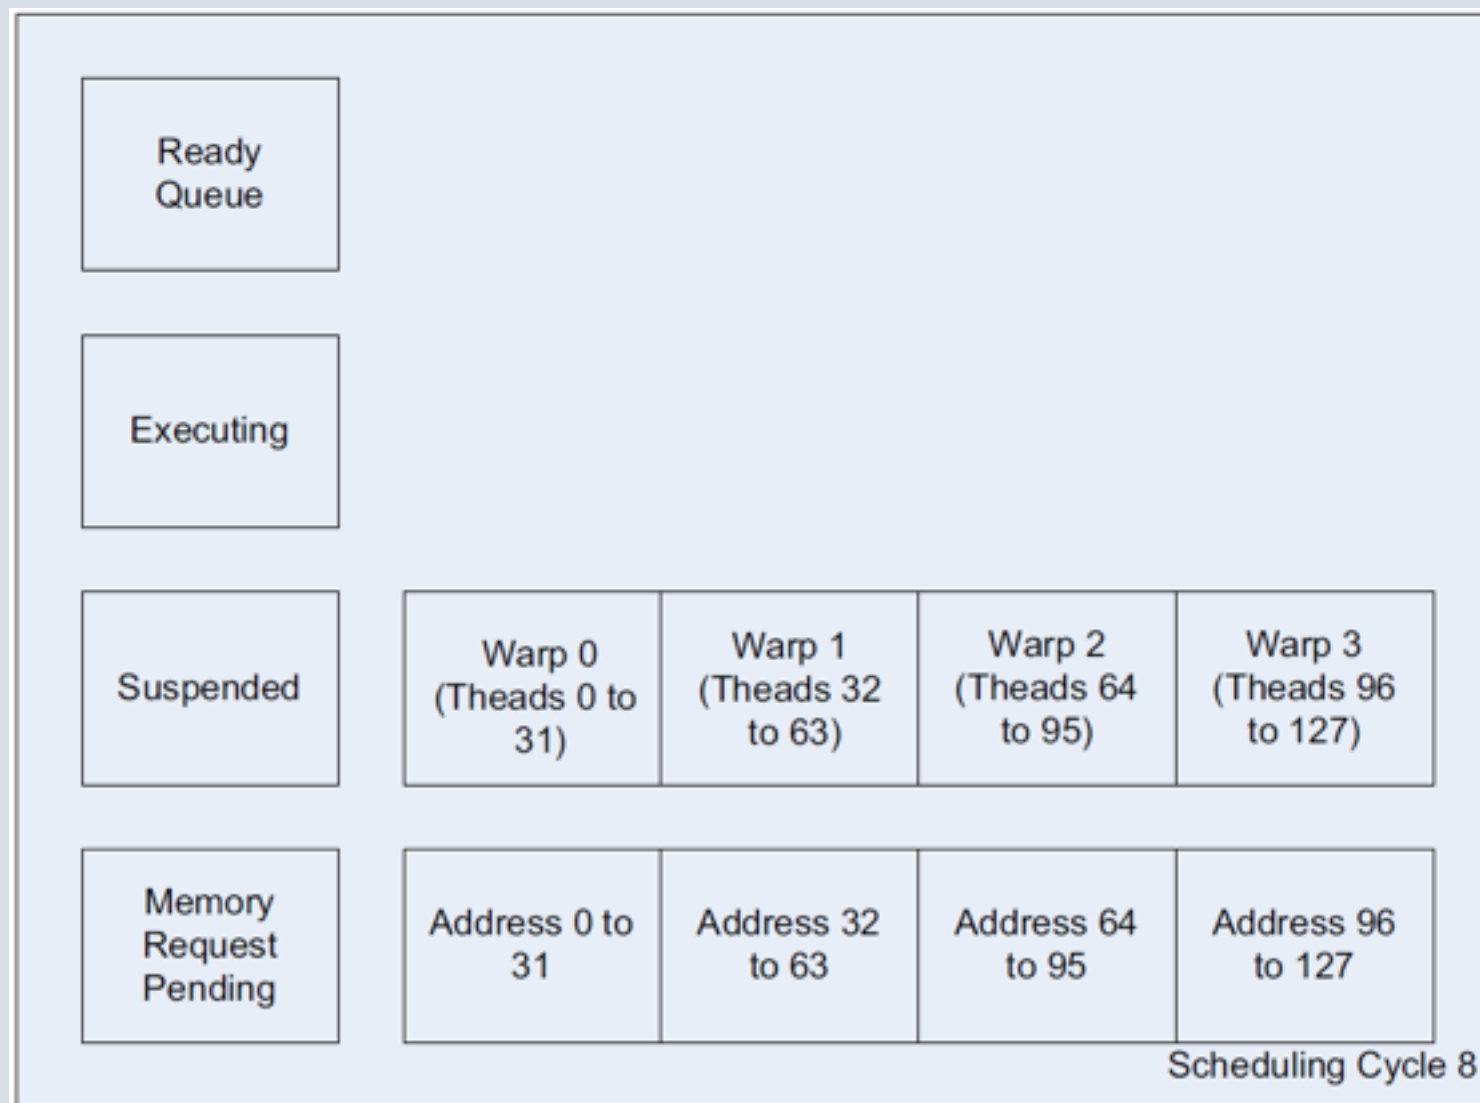
- ✓ Execute in warps of 32 threads



# CUDA Scheduling

## ●Block Scheduling

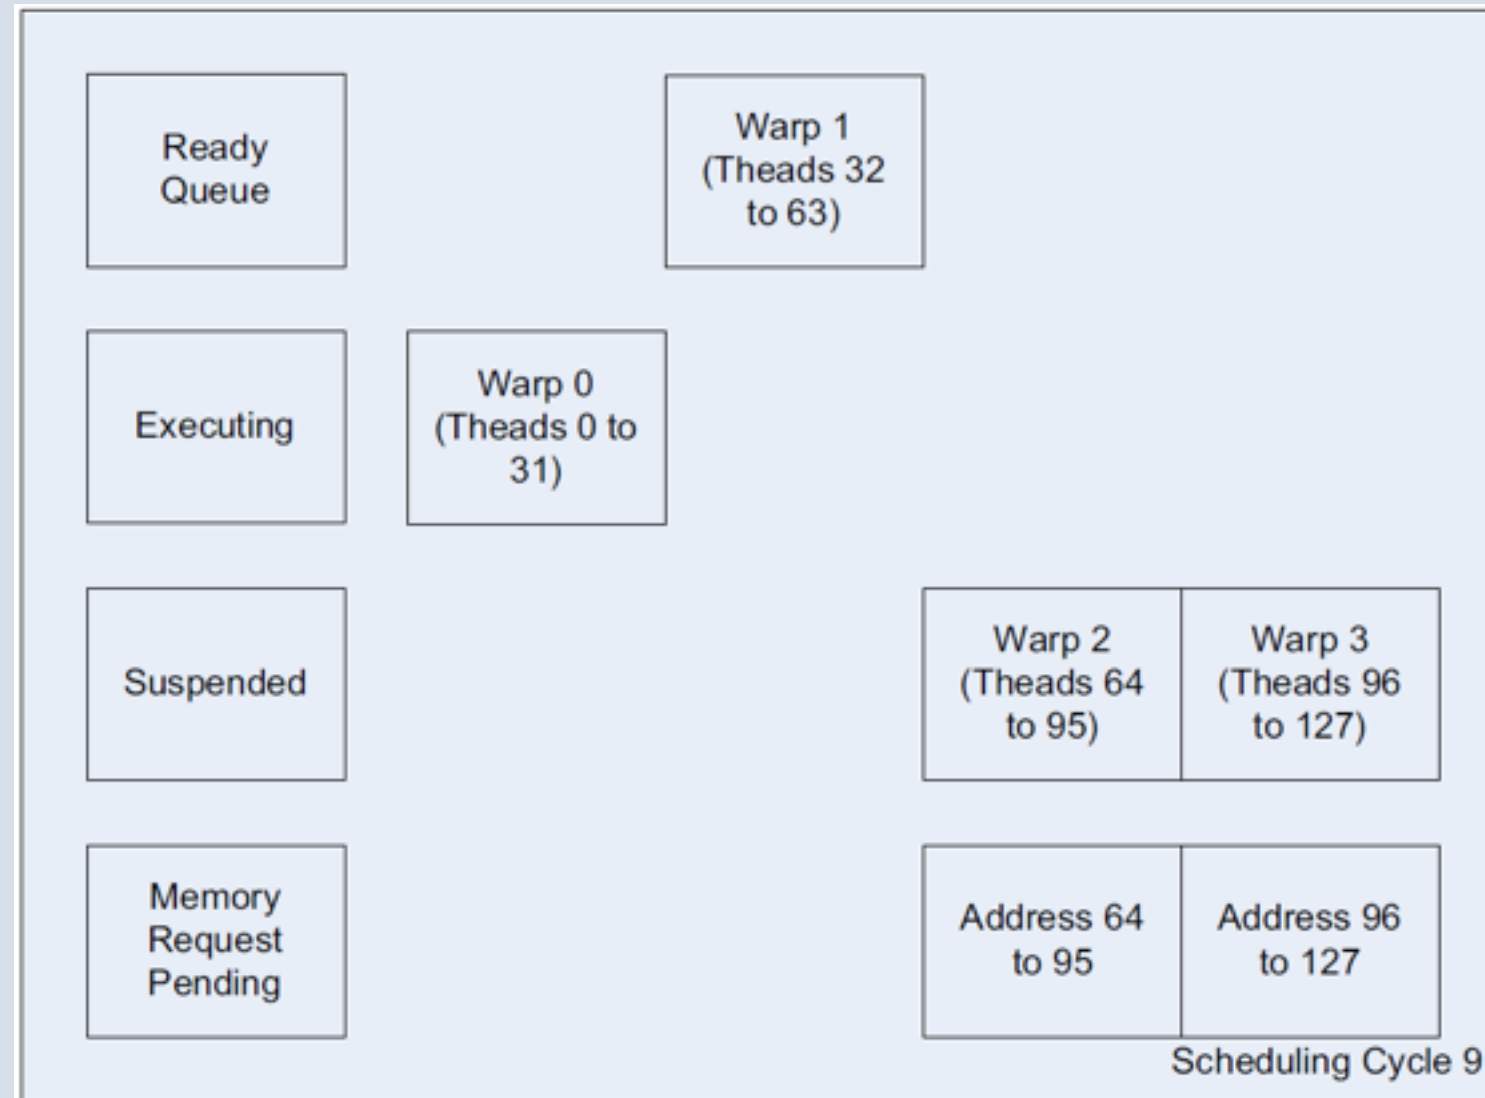
- ✓ Execute in warps of 32 threads



# CUDA Scheduling

## ●Block Scheduling

- ✓ Execute in warps of 32 threads





# Lab 2.1 Warp Scheduling

## ●理解线程束的调度机制

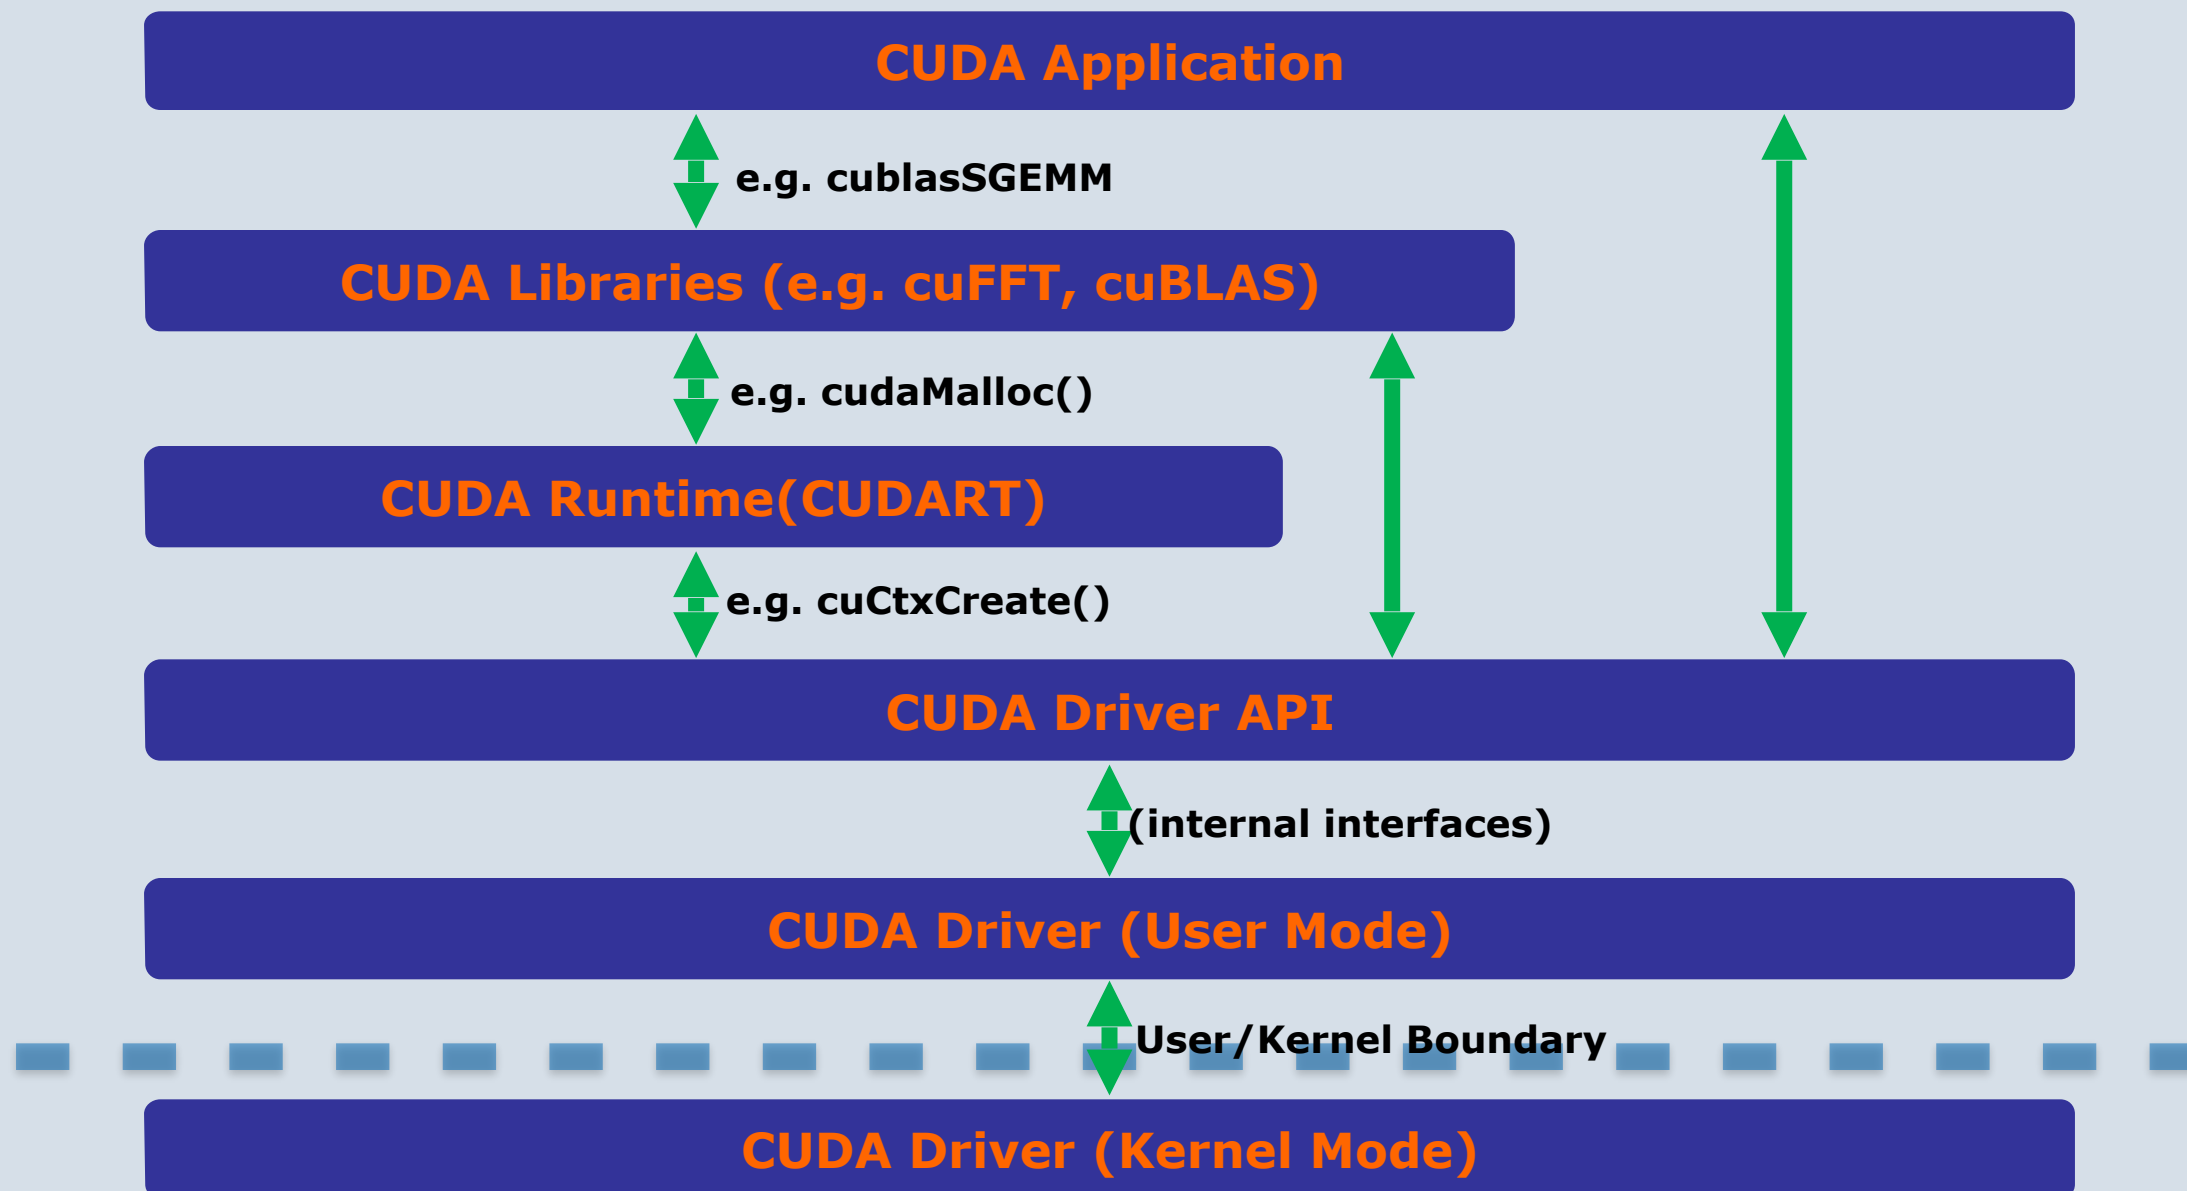
- ✓ 验证warp的线程数量
- ✓ 加入计时功能，对warp的调度时间进行输出，并绘出散点图进行分析
- ✓ 变大block和grid的大小会如何?
- ✓ 给出对线程束调度机制的理解
- ✓ 参见COOK 5.3 和WILT 7.3.3

```
__global__ void what_is_my_id(unsigned int * const block,
    unsigned int * const thread,
    unsigned int * const warp,
    unsigned int * const calc_thread)
{
    /* Thread id is block index * block size + thread offset into the block */
    const unsigned int thread_idx = (blockIdx.x * blockDim.x) + threadIdx.x;
    block[thread_idx] = blockIdx.x;
    thread[thread_idx] = threadIdx.x;
    /* Calculate warp using built in variable warpSize */
    warp[thread_idx] = threadIdx.x / warpSize;
    calc_thread[thread_idx] = thread_idx;
}
```

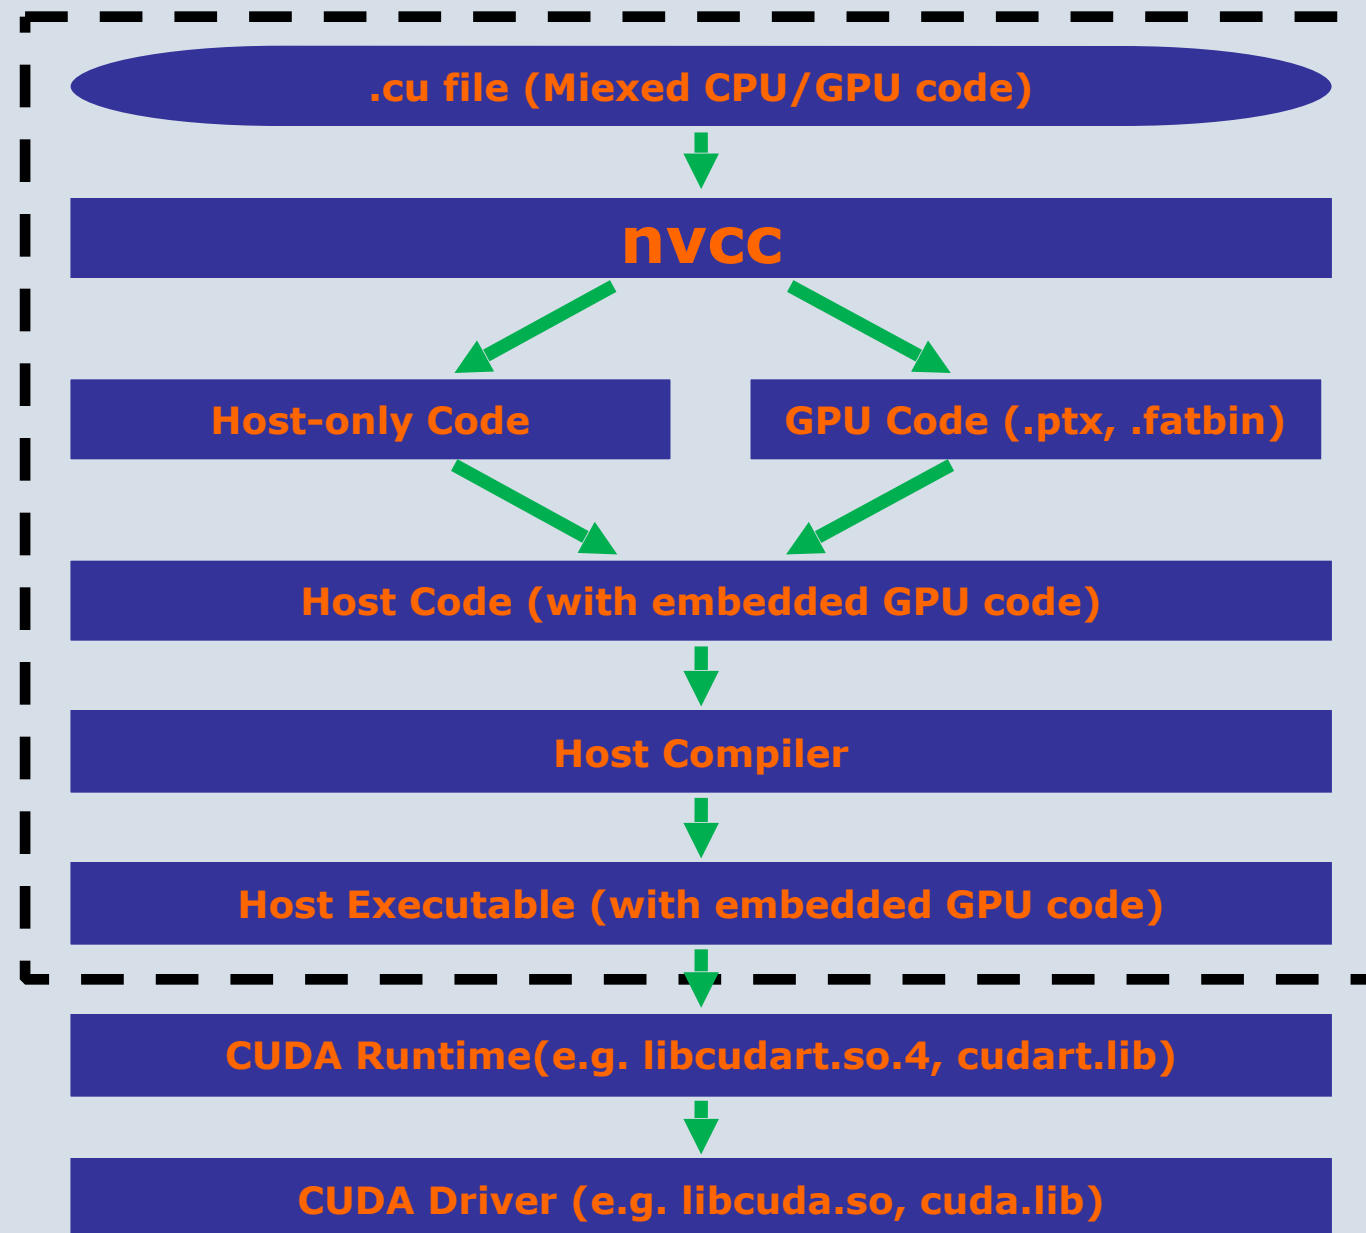
# Outline

- 1 **Multithreading**
- 2 **CUDA Abstraction**
- 3 **Kernel Execution**
- 4 **Warp Scheduling**
- 5 **CUDA Toolchain**

# Software Architecture



# Compiling Process



Refer to: "CUDA Compiler Driver NVCC", v7.5, 2015

Tonghua Su, School of Software, Harbin Institute of Technology, China

# NVCC Case Study

## ● Useful Options

- ✓ -g
- ✓ -G
- ✓ -lineinfo
- ✓ -o outfile
- ✓ -include xxx.h
- ✓ -l yyy.lib
- ✓ -arch sm\_??

# CUDA Software Ecosystem

- Numerical Analysis Tools

- ✓ e.g. Matlab, ArrayFire, Mathematica, LabView, Jacket

- GPU-Accelerated Libraries

- ✓ e.g. Thrust, cuDNN, cuFFT, cuBLAS

- Language and APIs

- ✓ e.g. CUDA Toolkit, OpenACC, CUDA Fortran, PyCUDA, PGI compilers

- Performance Analysis Tools

- ✓ e.g. Nsight, NVIDIA Visual Profiler,

- Debugging Solutions

- ✓ e.g. Nsight, CUDA-GDB, CUDA-MEMCHECK

# Quiz

- 请在命令行使用**nvcc**编译你的**cuda**程序
  - ✓ 在**vs2013**里，**nvcc**的调用在哪里，尝试解释其中的参数含义